



King Abdullah University
of Science and Technology



GMRFLib 2.0 (?)

Henrique Ap. Laureano

`henrique.laureano@kaust.edu.sa` \wedge `http://mynameislaure.github.io/`

`/KAUST/CEMSE/STAT`

Summer Semester
2018

Contents

| | |
|--|----|
| Markov Random Fields | 2 |
| Gaussian Markov Random Fields | 2 |
| Basic Properties | 4 |
| Conditional Properties | 4 |
| Markov Properties | 5 |
| Conditional Specification | 6 |
| Multivariate GMRFs | 8 |
| Exact Algorithms for GMRFs | 10 |
| Why are Exact Algorithms Important? | 10 |
| Some Basic Linear Algebra | 10 |
| Sampling from a GMRF | 11 |
| Sampling from a GMRF Conditioned on Linear Constraints | 11 |
| The Cholesky Factorization of Q | 13 |
| Interpretation of the Cholesky Triangle | 13 |
| Cholesky Factorization of Band Matrices | 15 |
| Reordering Techniques: Band Matrices | 15 |
| Reordering Techniques: General Sparse Matrices | 16 |

Markov Random Fields

"Statistical modeling of a finite collection of spatial random variables is often done through a Markov random field (MRF). A MRF is specified through the set of conditional distributions of one component given all the others. This enables one to focus on a single random variable at a time and leads to simple computational procedures for simulating MRFs, in particular for Bayesian inference via Markov chain Monte Carlo (MCMC)." In the Gaussian (*aka* Normal) case, we have the so called Gaussian MRFs (GMRFs).

Gaussian Markov Random Fields

"A GMRF is simply a Gaussian distributed random vector \mathbf{x} , which obeys some conditional independence properties. That is, for some $i \neq j$, then

$$x_i \perp x_j \mid \mathbf{x}_{-\{i,j\}}, \quad (1)$$

meaning that conditioned on $\mathbf{x}_{-\{i,j\}}$, x_i and x_j are independent. This conditional independence is represented using an (undirected) labeled graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$, where $\mathcal{V} = \{1, \dots, n\}$ is the set of vertices, an $\mathcal{E} = \{\{i, j\} : i, j \in \mathcal{V}\}$ is the set of edges in the graph. For all $i, j \in \mathcal{V}$, the edge $\{i, j\}$ is not included in \mathcal{E} if (1) holds, and included otherwise. Figure 1 displays such a graph, where $n = 4$ and $\mathcal{E} = \{\{1, 2\}, \{2, 3\}, \{3, 4\}, \{4, 1\}\}$. From this graph we deduce that $x_2 \perp x_4 \mid \mathbf{x}_{\{1,3\}}$ and $x_1 \perp x_3 \mid \mathbf{x}_{\{2,4\}}$. A central goal is now to specify a GMRF \mathbf{x} with conditional independence properties in agreement with some given graph \mathcal{G} ."




Figure 1: A conditional independence graph.

Theorem 1.

Let \mathbf{x} be Gaussian distributed with a symmetric and positive definite (SPD) precision matrix \mathbf{Q} , then for $i \neq j$

$$x_i \perp x_j \mid \mathbf{x}_{-\{i,j\}} \iff Q_{i,j} = 0.$$

"So any SPD precision matrix \mathbf{Q} with $Q_{2,4} = Q_{4,2} = Q_{1,3} = Q_{3,1} = 0$ has conditional independence properties as displayed in Figure 1." A precision matrix \mathbf{Q} that may correspond to this is presented in (2). "We then say that \mathbf{x} is a GMRF with respect to \mathcal{G} . A formal definition follows."

$$\mathbf{Q} = \begin{matrix} & \begin{matrix} 1 & 2 & 3 & 4 \end{matrix} \\ \begin{matrix} 1 \\ 2 \\ 3 \\ 4 \end{matrix} & \begin{pmatrix} Q_{1,1} & Q_{1,2} & 0 & Q_{1,4} \\ Q_{2,1} & Q_{2,2} & Q_{2,3} & 0 \\ 0 & Q_{3,2} & Q_{3,3} & Q_{3,4} \\ Q_{4,1} & 0 & Q_{4,3} & Q_{4,4} \end{pmatrix} \end{matrix}. \quad (2)$$

Definition 1 (GMRF).

A random vector $\mathbf{x} = (x_1, \dots, x_n)^\top \in \mathbb{R}^n$ is called a GMRF wrt the labeled graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ with mean $\boldsymbol{\mu}$ and SPD precision matrix \mathbf{Q} , iff its density has the form

$$\pi(\mathbf{x}) = (2\pi)^{-n/2} |\mathbf{Q}|^{1/2} \exp \left(-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^\top \mathbf{Q} (\mathbf{x} - \boldsymbol{\mu}) \right) \quad (3)$$

and

$$Q_{i,j} \neq 0 \iff \{i, j\} \in \mathcal{E} \quad \forall \quad i \neq j.$$

"The case where \mathbf{Q} is singular still provides a GMRF with an explicit form for its joint density, but the joint density is improper. Such specifications cannot be used as data models, but can be used as priors as long as they yield proper posteriors. Here is a simple example of a (proper) GMRF."

Example 1.

"Let $\{x_t\}$ be a stationary autoregressive process of order one, *i.e.*, $x_t | x_{t-1} = \phi x_{t-1} + \epsilon_t$, for $t = 2, \dots, n$, where $|\phi| < 1$ and ϵ_t are independent normally distributed zero mean innovations with unit variance. Further assume that x_1 is normal with mean zero and variance $1/(1\phi^2)$, which is simply the stationary distribution of this process. Then \mathbf{x} is a GMRF wrt to \mathcal{G} where $\mathcal{E} = \{\{1, 2\}, \{2, 3\}, \dots, \{n-1, n\}\}$. The precision matrix has nonzero elements $Q_{i,j} = -\phi$ for $|i-j| = 1$, $Q_{1,1} = Q_{n,n} = 1$ and $Q_{i,i} = 1 + \phi^2$ for $i = 2, \dots, n-1$." Considering, *e.g.*, $n = 5$, the precision matrix \mathbf{Q} is given by

$$\begin{matrix} & \begin{matrix} 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{pmatrix} 1 & -\phi & 0 & 0 & 0 \\ -\phi & 1+\phi^2 & -\phi & 0 & 0 \\ 0 & -\phi & 1+\phi^2 & -\phi & 0 \\ 0 & 0 & -\phi & 1+\phi^2 & -\phi \\ 0 & 0 & 0 & -\phi & 1 \end{pmatrix} \end{matrix}.$$

"This example nicely illustrates why GMRFs are so useful." "Only $n+2(n-1) = 3n-2$ of the n^2 terms in \mathbf{Q} are nonzero. The sparse precision matrix makes fast $\mathcal{O}(n)$ algorithms for the simulation of autoregressive processes possible."

Basic Properties

Conditional Properties

"Although a GMRF can be seen as a general multivariate Gaussian random variable, some properties simplify and some characteristics are easier to compute. For example, conditional distributions are easier to compute due to the sparse precision matrix. To see this, we split \mathcal{V} into the nonempty sets A and $B = \mathcal{V} \setminus A$. Partition \mathbf{x} , $\boldsymbol{\mu}$ and \mathbf{Q} accordingly, *i.e.*,

$$\mathbf{x} = \begin{pmatrix} \mathbf{x}_A \\ \mathbf{x}_B \end{pmatrix}, \quad \boldsymbol{\mu} = \begin{pmatrix} \boldsymbol{\mu}_A \\ \boldsymbol{\mu}_B \end{pmatrix} \quad \text{and} \quad \mathbf{Q} = \begin{pmatrix} \mathbf{Q}_{AA} & \mathbf{Q}_{AB} \\ \mathbf{Q}_{BA} & \mathbf{Q}_{BB} \end{pmatrix}.$$

We also need the notion of a subgraph \mathcal{G}^A , which is the graph restricted to A : the graph we obtain after removing all nodes not belonging to A and all edges where at least one node does not belong to A . Then the following theorem holds."

Theorem 2.

Let \mathbf{x} be a GMRF wrt \mathcal{G} with mean $\boldsymbol{\mu}$ and SPD precision matrix \mathbf{Q} . Let $A \subset \mathcal{V}$ and $B = \mathcal{V} \setminus A$ where $A, B \neq \emptyset$. The conditional distribution of $\mathbf{x}_A | \mathbf{x}_B$ is then a GMRF wrt the subgraph \mathcal{G}^A with mean $\boldsymbol{\mu}_{A|B}$ and SPD precision matrix $\mathbf{Q}_{A|B}$, where

$$\mu_{A|B} = \mu_A - Q_{AA}^{-1} Q_{AB} (x_B - \mu_B) \quad \text{and} \quad Q_{A|B} = Q_{AA}.$$

"The expression for the conditional mean $\mu_{A|B}$ involves the inverse Q_{AA}^{-1} , but only in a way such that we can write $\mu_{A|B} = \mu_A - b$, where b is the solution of a sparse linear system $Q_{AA} b = Q_{AB} (x_B - \mu_B)$. Note that the term Q_{AB} is nonzero only for those vertices in A that have an edge to a vertex in B , so usually only a few terms will enter in this matrix-vector product. In the special case $A = \{i\}$, the expressions simplify to

$$\mu_{i|-i} = \mu_i - \sum_{j:j \sim i} \frac{Q_{i,j}}{Q_{i,i}} (x_j - \mu_j) \quad \text{and} \quad Q_{i|-i} = Q_{i,i}. \quad (4)$$

Here we used the notation $j : j \sim i$ to indicate a sum over all vertices j that are neighbors to vertex i , i.e., $\{i, j\} \in \mathcal{E}$. So $Q_{i,i}$ is the conditional precision of x_i and the conditional expectation of x_i is a weighted mean of neighboring x_j s with weights $-Q_{i,j}/Q_{i,i}$."

Example 2.

"We continue with Example 1. From (4) we obtain the conditional mean and precision of $x_i \mid \mathbf{x}_{-i}$,"

$$\begin{aligned} \mu_{i|-i} &= 0 - \left[\frac{-\phi}{1 + \phi^2} (x_{i-1} - 0) + \frac{-\phi}{1 + \phi^2} (x_{i+1} - 0) \right] \\ &= 0 - \frac{-\phi}{1 + \phi^2} (x_{i-1} + x_{i+1}) \\ &= \frac{\phi}{1 + \phi^2} (x_{i-1} + x_{i+1}), \quad \text{and} \quad Q_{i|-i} = 1 + \phi^2, \quad 1 < i < n. \end{aligned}$$

Markov Properties

"The graph \mathcal{G} of a GMRF is defined through looking at which x_i and x_j are conditionally independent, the so-called *pairwise* Markov property. However, more general Markov properties can be derived from \mathcal{G} .

A *path* from vertex i_1 to vertex i_m is a sequence of distinct nodes in \mathcal{V} , i_1, i_2, \dots, i_m , for which $(i_j, i_{j+1}) \in \mathcal{E}$ for $j = 1, \dots, m-1$. A subset $C \subset \mathcal{V}$ *separates* two nodes $i \notin C$ and $j \notin C$, if every path from i to j contains at least one node from C . Two disjoint sets $A \subset \mathcal{V} \setminus C$ and $B \subset \mathcal{V} \setminus C$ are separated by C , if all $i \in A$ and $j \in B$ are separated by C . In other words, we cannot walk on the graph starting somewhere in A ending somewhere in B without passing through C . The global Markov property, is that

$$\mathbf{x}_A \perp \mathbf{x}_B \mid \mathbf{x}_C$$

for all mutually disjoint sets A , B and C where C separates A and B , and A and B are nonempty."

Theorem 3.

Let \mathbf{x} be a GMRF wrt \mathcal{G} , then \mathbf{x} obeys the global Markov property.

Conditional Specification

"It is common to specify a GMRF implicitly through the so-called full conditionals $\{\pi(x_i \mid \mathbf{x}_{-i})\}$." "However, the full conditionals cannot be specified completely arbitrarily, as we must ensure that they correspond to a proper joint density.

A conditional specification defines the full conditional $\{\pi(x_i \mid \mathbf{x}_{-i})\}$ as normal with moments

$$\mathbb{E}(x_i \mid \mathbf{x}_{-i}) = \mu_i + \sum_{j \neq i} \beta_{i,j}(x_j - \mu_j) \quad \text{and} \quad \text{Precision}(x_i \mid \mathbf{x}_{-i}) = \kappa_i > 0. \quad (5)$$

The rationale for such an approach, is that it is easier to specify the full conditionals than the joint distribution. Comparing (5) with (3), we can choose $\boldsymbol{\mu}$ as the mean, $Q_{i,i} = \kappa_i$, $\beta_{i,j} = -Q_{i,j}/Q_{i,i}$ to obtain the same full conditionals. However, since \mathbf{Q} is symmetric, we must require that

$$\kappa_i \beta_{i,j} = \kappa_j \beta_{j,i} \quad (6)$$

for all $i \neq j$." "In addition to the symmetry constraint (6), there is a joint requirement that \mathbf{Q} is SPD. Unfortunately, this is a *joint* property, which is hard to validate locally. One convenient approach that avoids this problem is to choose a diagonally dominant parametrization that ensures \mathbf{Q} to be SPD: $Q_{i,i} > \sum_j |Q_{i,j}|$ for all i . This implies that"

$$\sum_j |\beta_{i,j}| < 1, \quad \forall i.$$

Fixing $\boldsymbol{\mu} = \mathbf{0}$, using the full conditionals in Equation (5), and considering the symmetry constraint (6), the density of \mathbf{x} can then be expressed as

$$\begin{aligned} \log \pi(\mathbf{x}) &= \text{const} + \frac{1}{2} \log |\mathbf{Q}| - \frac{1}{2} \mathbf{x}^\top \mathbf{Q} \mathbf{x} \quad (\text{starting from (3)}) \\ &= \text{const} + \frac{1}{2} \sum_{i \neq j} Q_{i,j} x_i x_j - \frac{1}{2} \sum_{i=1}^n Q_{i,i} x_i^2 \\ &= \text{const} - \frac{1}{2} \sum_{i \neq j} \kappa_i \beta_{i,j} x_i x_j - \frac{1}{2} \sum_{i=1}^n \kappa_i x_i^2; \end{aligned}$$

"hence, \mathbf{x} is zero mean GMRF provided \mathbf{Q} is SPD."

Example 3.

"The image in Figure 2(a) is a 256×256 gamma camera image of a phantom designed to reflect structure expected from cancerous bones. Each pixel in the circular part of the image, \mathcal{I} , represent the gamma radiation count, where a black pixel represents (essentially) zero counts and a white pixel the maximum count. The image is quite noisy and the task in this example is to (try to) remove the noise. The noise process is quite accurately described by a Poisson distribution, so that for each pixel i , the recorded count y_i relates to the true signal η_i , as $y_i \sim \text{Poisson}(\eta_i)$. For simplicity, we will use the approximation that"




Figure 2: Panel (a) shows the raw x-ray image, (b) shows the square-root transformed image, and (c) shows the restored image (the posterior mean).

$$\sqrt{y_i} \mid \eta_i \sim \mathcal{N}\left(\sqrt{\eta_i}, \frac{1}{4}\right), \quad i \in \mathcal{I}$$

"and the square-root transformed image is displayed in Figure 2(b). Taking a Bayesian approach, we need to specify a prior distribution for the (square-root-transformed) image $\mathbf{x} = (x_1, \dots, x_n)^\top$, where $x_i = \sqrt{\eta_i}$. (We need η_i to be (somewhat) larger than zero for this approximation to be adequate.) Although this is a daunting problem in general, for such noise-removal tasks it is usually sufficient to specify the prior to be informative for how the true image behaves locally. Since the image itself is locally smooth, we might specify the prior through the full conditionals (5). Using the full conditionals we only need to answer questions like: *What if we do not know the true signal in pixel i , but all others; what is then our belief in x_i ?* One choice, is to set $\beta_{i,j}$ to zero unless j is one of the four nearest neighbors of i ; $N_4(i)$, say. As we have no particular preference for direction, we might take for each i ,

$$\beta_{i,j} = \frac{\delta}{4}, \quad j \in N_4(i)$$

where δ is considered as fixed. Further, we take κ_i to be common (and unknown) for all i , and restrict δ to $|\delta| < 1$ so that the (prior) precision matrix is diagonally dominant. (We ignore here some corrections at the boundary where a boundary pixels may have less than four neighbors.) We take further $\boldsymbol{\mu} = \mathbf{0}$ and a (conjugate) $\Gamma(a, b)$ prior for κ (with density $\propto \kappa^{a-1} \exp(-b\kappa)$), and then the posterior for (\mathbf{x}, κ) reads

$$\begin{aligned} \pi(\mathbf{x}, \kappa \mid \mathbf{y}) &\propto \pi(\mathbf{x}, \kappa) \pi(\kappa) \prod_{i \in \mathcal{I}} \pi(y_i \mid x_i) \\ &\propto \kappa^{a-1} \exp(-b\kappa) |\mathbf{Q}_{prior}(\kappa)|^{1/2} \exp\left(-\frac{1}{2} \mathbf{x}^\top \mathbf{Q}_{post}(\kappa) \mathbf{x} + \mathbf{b}^\top \mathbf{x}\right). \end{aligned}$$

Here, $b_i = 4\sqrt{y_i}$ for $i \in \mathcal{I}$ and zero otherwise, $\mathbf{Q}_{post}(\kappa) = \mathbf{Q}_{prior}(\kappa) + \mathbf{D}$ where \mathbf{D} is a diagonal matrix where $D_{i,i} = 4$ if $i \in \mathcal{I}$ and zero otherwise, and




Figure 3: Panels (a) to (e) show five consecutive frames of a three-dimensional confocal microscopy image.

$$\mathbf{Q}_{prior}(\kappa)_{i,j} = \kappa \begin{cases} 1, & i = j \\ \delta/4, & j \in N_4(i) \\ 0, & \text{otherwise.} \end{cases}$$

Conditioned on κ and the observations, then \mathbf{x} is a GMRF with precision matrix \mathbf{Q}_{post} and where the mean $\boldsymbol{\mu}_{post}$ is given by the solution of"

$$\mathbf{Q}_{post}\boldsymbol{\mu}_{post} = \mathbf{b}.$$

Multivariate GMRFs

"To fix ideas, we will consider a generalization of Example 3 where the observations are now sequences of images. The sequence can either be a movie where each frame in the sequence is indexed by time, or the height where recorded a three-dimensional object as a set of two-dimensional images. Other examples include a temporal version of spatial models of disease counts in each administrative region of a country. Figure 3 shows five consecutive frames of three-dimensional cells taken by confocal microscopy. The first frame has a lot of noise, but the signal gets stronger farther up in the image stack. We consider the same problem as for Example 3; we want to estimate the true signal in the presence of the noise. The five frames represent the same three-dimensional object, but at different height.

We can use this information when we specify the full conditionals. It is then both easier and more natural to specify a multivariate version of the full conditionals (5), which we now will describe. Let \mathbf{x}_i represent all the $p = 5$ observations at pixel i

$$\mathbf{x}_i = (x_{i,1}, x_{i,2}, x_{i,3}, x_{i,4}, x_{i,5})^\top.$$

Here, $x_{i,2}$ is the pixel at location i in frame 2 and so on. The conditional specification (5) extends naturally to

$$\mathbb{E}(x_i | x_{-i}) = \mu_i - \sum_{j:j \sim i} \beta_{i,j}(x_j - \mu_j) \quad \text{and} \quad \text{Precision}(x_i | x_{-i}) = \kappa_i > 0, \quad (7)$$

for some $p \times p$ matrices $\{\beta_{i,j}\}$ and $\{\kappa_i\}$. In this formulation, we can now specify that our knowledge of $x_{i,3}$ might benefit of knowing $x_{i,2}$ and $x_{i,4}$. These pixels are in the same x_i vector, although they represent the i th pixel at the previous and next frame. Additionally, we can have dependency from neighbors within the same frame, such as $\{x_{j,3}, j \in N_4(i)\}$. In short, we can specify how x_i depends on $\{x_j\}, j \neq i$, and thinking about neighbors that are (small p -) vectors.

The conditional specification in this example motivates the introduction of a multivariate GMRF, which we denote as MGMRF $_p$. Its definition is a direct extension of (1). Let $\mathbf{x} = (x_1^\top, \dots, x_n^\top)^\top$ be Gaussian distributed, where each x_1 is a p -vector. Similarly, let $\boldsymbol{\mu} = (\mu_1^\top, \dots, \mu_n^\top)^\top$ denote the mean and $\tilde{\mathbf{Q}} = (\tilde{Q}_{i,j})$ the precision matrix with $p \times p$ elements $\tilde{Q}_{i,j}$.

Definition 2 (MGMRF $_p$).

A random vector $\mathbf{x} = (x_1^\top, \dots, x_n^\top)^\top$ where $\dim(x_i) = p$, is called a MGMRF $_p$ wrt $\mathcal{G} = (\mathcal{V} = \{1, \dots, n\}, \mathcal{E})$ with mean $\boldsymbol{\mu}$ and SPD precision matrix $\tilde{\mathbf{Q}}$, iff its density has the form

$$\begin{aligned} \pi(\mathbf{x}) &= (2\pi)^{-np/2} |\tilde{\mathbf{Q}}|^{1/2} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^\top \tilde{\mathbf{Q}}(\mathbf{x} - \boldsymbol{\mu})\right) \\ &= (2\pi)^{-np/2} |\tilde{\mathbf{Q}}|^{1/2} \exp\left(-\frac{1}{2} \sum_{ij} (x_i - \mu_i)^\top \tilde{Q}_{i,j} (x_j - \mu_j)\right) \end{aligned}$$

and

$$\tilde{Q}_{i,j} \neq 0 \quad \Longleftrightarrow \quad \{i, j\} \in \mathcal{E} \quad \forall \quad i \neq j.$$

"It is important to note that a size n MGMRF $_p$ is just another GMRF of dimension np ; so all our previous results and forthcoming sparse matrix algorithms for GMRFs also apply for a MGMRF $_p$. However, some results have easier interpretation using the block formulation, such as

$$x_i \perp x_j \mid x_{-\{i,j\}} \quad \Longleftrightarrow \quad \tilde{Q}_{i,j} = 0$$

and

$$\mathbb{E}(x_i | x_{-i}) = \mu_i - \tilde{Q}_{i,j}^{-1} \sum_{j:j \sim i} \tilde{Q}_{i,j}(x_j - \mu_j) \quad \text{and} \quad \text{Precision}(x_i | x_{-i}) = \tilde{Q}_{i,i}. \quad (8)$$

From Equation (8), we can obtain the consistency requirements for the conditional specification Equation (7) by choosing

$$\tilde{\mathbf{Q}}_{i,j} = \begin{cases} \kappa_i \beta_{i,j}, & i \neq j \\ \kappa_i, & i = j. \end{cases}$$

Since $\tilde{\mathbf{Q}}_{i,j} = \tilde{\mathbf{Q}}_{j,i}^\top$, then we have the requirement that $\kappa_i \beta_{i,j} = \beta_{j,i}^\top \kappa_j$ for $i \neq j$, additionally to $\kappa_i > 0 \quad \forall i$. Finally, there is also the "global" requirement that $\tilde{\mathbf{Q}}$ must be SPD, which is equivalent to $(\mathbf{I} + (\beta_{i,j}))$ being SPD."

Exact Algorithms for GMRFs

GMRFs have "a nice connection with very efficient numerical algorithms for sparse matrices. This connection allows for exact algorithms for GMRFs. We will now discuss this connection, starting with various exact algorithms to efficiently sample from GMRFs. This includes solving tasks like unconditional and conditional sampling, sampling under linear hard and soft constraints, evaluating the log-density of a (possibly constrained) GMRF at a particular value, and computing marginal variances for (possibly constrained) GMRFs. Although all these tasks are formally "just matrix algebra," we need to ensure that we take advantage of the sparse precision matrix \mathbf{Q} in all steps so computations can make use of the efficient numerical algorithms for sparse matrices developed in the computational sciences literature. Further, we can derive all the algorithms for sparse matrices by considering conditional independence properties of GMRFs. The core of all algorithms is the Cholesky factorization $\mathbf{Q} = \mathbf{L}\mathbf{L}^\top$ of the precision matrix \mathbf{Q} , where \mathbf{L} is a lower-triangular matrix."

Why are Exact Algorithms Important?

"Exact efficient algorithms are generally preferable when they exist, even though they apparently require algorithms that are more involved than simple iterative ones. Computational feasibility is important even for statistical modeling, as a statistical model is not of much use if we cannot do inference efficiently enough to satisfy the end-user.

Sampling from a GMRF can be done exactly using the Cholesky factorization of the precision matrix". "In the spatial case, it turns out that we can (typically) sample a GMRF exactly at the cost of $\mathcal{O}(n^{3/2})$ operations". "The exact algorithm can further produce independent samples at the cost of $\mathcal{O}(n \log n)$ each."

Some Basic Linear Algebra

"Let \mathbf{A} be SPD, then there exists a unique (Cholesky) factorization $\mathbf{A} = \mathbf{L}\mathbf{L}^\top$, where \mathbf{L} is lower triangular and called the Cholesky triangle. This factorization is the starting point for solving $\mathbf{A}\mathbf{y} = \mathbf{b}$ by first solving $\mathbf{L}\mathbf{v} = \mathbf{b}$ and then $\mathbf{L}^\top \mathbf{y} = \mathbf{v}$. The first linear system $\mathbf{L}\mathbf{v} = \mathbf{b}$ is solved directly using forward substitution

$$v_i = \frac{1}{L_{i,i}} \left(b_i - \sum_{j=1}^{i-1} L_{i,j} v_j \right), \quad i = 1, \dots, n,$$

whereas $\mathbf{L}^\top \mathbf{y} = \mathbf{v}$ is solved using backward substitution

$$y_i = \frac{1}{L_{i,i}} \left(v_i - \sum_{j=1+1}^n L_{j,i} y_j \right), \quad i = n, \dots, 1.$$

Computing $\mathbf{A}^{-1} \mathbf{Y}$, where \mathbf{Y} is a $n \times k$ matrix, is done by computing $\mathbf{A}^{-1} \mathbf{Y}_j$ for each of the k columns \mathbf{Y}_j using the algorithm above. Note that \mathbf{A} needs to be factorized only once. Note that in the case $\mathbf{Y} = \mathbf{I}$ (and $k = n$), the inverse of \mathbf{A} is computed".

Sampling from a GMRF

"Sampling from a GMRF can be done using the following steps: sample $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$, *i.e.*, n standard normal variables, solve $\mathbf{L}^\top \mathbf{v} = \mathbf{z}$, and compute $\mathbf{x} = \boldsymbol{\mu} + \mathbf{v}$. The sample \mathbf{x} has the correct distribution as $\mathbb{E}(\mathbf{v}) = \mathbf{0}$ and $\text{Cov}(\mathbf{v}) = \mathbf{L}^{-\top} \mathbf{I} \mathbf{L}^{-1} = (\mathbf{L} \mathbf{L}^\top)^{-1} = \mathbf{Q}^{-1}$. The log-density of \mathbf{x} ,

$$\log \pi(\mathbf{x}) = -\frac{n}{2} \log 2\pi + \frac{1}{2} \log |\mathbf{Q}| - \frac{1}{2} \underbrace{(\mathbf{x} - \boldsymbol{\mu})^\top \mathbf{Q} (\mathbf{x} - \boldsymbol{\mu})}_q$$

is evaluated as follows. If \mathbf{x} is sampled using the algorithm above, then $q = \mathbf{z}^\top \mathbf{z}$, otherwise, we compute $\mathbf{w} = \mathbf{x} - \boldsymbol{\mu}$, $\mathbf{u} = \mathbf{Q} \mathbf{w}$ and then $q = \mathbf{w}^\top \mathbf{u}$. Note that $\mathbf{u} = \mathbf{Q} \mathbf{w}$ is a sparse-matrix vector product, which can be computed efficiently:

$$u_i = Q_{i,i} w_i + \sum_{j:j \sim i} Q_{i,j} w_j,$$

where the diagonal term is added explicitly since i is not a neighbor of i . The determinant of \mathbf{Q} can be found from the Cholesky factorization: $|\mathbf{Q}| = |\mathbf{L} \mathbf{L}^\top| = |\mathbf{L}|^2$. Since \mathbf{L} is lower triangular, we obtain

$$\frac{1}{2} \log |\mathbf{Q}| = \sum_i \log L_{i,i}.$$

Conditional sampling of \mathbf{x}_A conditioned on \mathbf{x}_B , as described in Theorem 2, is similar: factorize \mathbf{Q}_{AA} and *solve*

$$\mathbf{Q}_{AA}(\boldsymbol{\mu}_{A|B} - \boldsymbol{\mu}_A) = -\mathbf{Q}_{AB}(\mathbf{x}_B - \boldsymbol{\mu}_B)$$

for $\mathbf{x}_{A|B}$, using forward and backward substitution. The (sparse) matrix vector product on the right-hand side is computed using only the nonzero terms in \mathbf{Q}_{AB} . The remaining steps are the same."

Sampling from a GMRF Conditioned on Linear Constraints

"In practical applications, we often want to sample from a GMRF under a linear constraint,

$$\mathbf{A} \mathbf{x} = \mathbf{e},$$

for a $k \times n$ matrix \mathbf{A} of rank k . The common case is that $k \ll n$. A brute-force approach is to directly compute the conditional (Gaussian) density $\pi(\mathbf{x} \mid \mathbf{Ax} = \mathbf{e})$, but this will reveal that the corresponding precision matrix is (usually) not sparse anymore. For example, if $x_i \perp x_j \mid \mathbf{x}_{-\{i,j\}}$ without the constraint, then a sum-to-zero constraint $\sum x_i = 0$ makes x_i and x_j negatively correlated. In order not to lose computational efficiency, we must approach this problem in a more subtle way by correcting an unconstrained sample \mathbf{x} to obtain a constrained sample \mathbf{x}^c :

$$\mathbf{x}^c = \mathbf{x} - \mathbf{Q}^{-1} \mathbf{A}^\top (\mathbf{AQ}^{-1} \mathbf{A}^\top)^{-1} (\mathbf{Ax} - \mathbf{e}). \quad (9)$$

A direct calculation shows that \mathbf{x}^c has the correct mean and covariance. A closer look at (9) makes it clear that all the matrix terms are easy to compute: $\mathbf{Q}^{-1} \mathbf{A}^\top$ just solves k linear systems of type $\mathbf{Q} \mathbf{v}_j = (\mathbf{A}^\top)_j$, for $j = 1, \dots, k$, whereas $\mathbf{AQ}^{-1} \mathbf{A}^\top$ is a $k \times k$ matrix and its Cholesky factorization is fast to compute since k is small in typical applications. Note that the extra cost of having k constraints is $\mathcal{O}(nk^2)$, hence, negligible when k is small.

Evaluating the constrained log density perhaps needs more attention, since $\mathbf{x} \mid \mathbf{Ax}$ is singular with rank $n - k$. However, the following identity can be used:

$$\pi(\mathbf{x} \mid \mathbf{Ax}) = \frac{\pi(\mathbf{Ax} \mid \mathbf{x}) \pi(\mathbf{x})}{\pi(\mathbf{Ax})}. \quad (10)$$

Note now that all terms on the right-hand side are easy to compute: $\pi(\mathbf{x})$ is a GMRF with mean $\boldsymbol{\mu}$ and precision matrix \mathbf{Q} , $\pi(\mathbf{Ax})$ is (k -dimensional) Gaussian with mean $\mathbf{A}\boldsymbol{\mu}$ and covariance $\mathbf{AQ}^{-1} \mathbf{A}^\top$, while $\pi(\mathbf{Ax} \mid \mathbf{x})$ is either 0 (when the configuration \mathbf{x} is inconsistent with the value of \mathbf{Ax} or equal to $|\mathbf{AA}^\top|^{-1/2}$."

Example 4.

"Let \mathbf{x} be n independent, zero-mean, normal random variables with variance $\{\sigma_i^2\}$. A sum-to-zero constrained sample \mathbf{x}^* can be generated from a sample of the unconstrained \mathbf{x} using

$$x_i^* = x_i - c \sigma_i^2, \quad \text{where} \quad c = \sum x_j / \sum \sigma_j^2, \quad i = 1, \dots, n.$$

The above construction can be generalized to condition on so-called *soft constraints*, which we condition on the k observations $\mathbf{y} = (y_1, \dots, y_k)^\top$, where

$$\mathbf{y} \mid \mathbf{x} \sim \mathcal{N}(\mathbf{Ax}, \boldsymbol{\Upsilon}).$$

Here, \mathbf{A} is a $k \times n$ matrix with rank k and $\boldsymbol{\Upsilon} > 0$. The conditional density for $\mathbf{x} \mid \mathbf{y}$ has precision matrix $\mathbf{Q} + \mathbf{A}^\top \boldsymbol{\Upsilon}^{-1} \mathbf{A}$, which is often a dense matrix. We can use the same approach as in Equation (9), which now generalizes to

$$\mathbf{x}^c = \mathbf{x} - \mathbf{Q}^{-1} \mathbf{A}^\top (\mathbf{AQ}^{-1} \mathbf{A}^\top + \boldsymbol{\Upsilon})^{-1} (\mathbf{Ax} - \boldsymbol{\epsilon}), \quad (11)$$

where $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{y}, \boldsymbol{\Upsilon})$. Similarly, if \mathbf{x} is a sample from $\pi(\mathbf{x})$ and then \mathbf{x}^c computed from Equation (11) is distributed as $\pi(\mathbf{x} \mid \mathbf{y})$. To evaluate the conditional density, we use the same approach as in Equation (10), which now reads

$$\pi(\mathbf{x} \mid \mathbf{y}) = \frac{\pi(\mathbf{y} \mid \mathbf{x})\pi(\mathbf{x})}{\pi(\mathbf{y})}.$$

Here, $\pi(\mathbf{x})$ is the density for a GMRF, $\pi(\mathbf{y} \mid \mathbf{x})$ is the density for a k -dimensional Gaussian with mean $\mathbf{A}\mathbf{x}$ and covariance matrix $\mathbf{\Upsilon}$, whereas $\pi(\mathbf{y})$ is the density for a k -dimensional Gaussian with mean $\mathbf{A}\boldsymbol{\mu}$ and covariance matrix $\mathbf{A}\mathbf{Q}^{-1}\mathbf{A}^\top + \mathbf{\Upsilon}$."

The Cholesky Factorization of \mathbf{Q}

"All the exact simulation algorithms are based on the Cholesky triangle \mathbf{L} , which is found by factorizing \mathbf{Q} into $\mathbf{L}\mathbf{L}^\top$. We will now discuss this factorization in more detail, show why sparse matrices allow for faster factorization, and how reordering the indices can speed up the computations."

Interpretation of the Cholesky Triangle

"The Cholesky factorization of \mathbf{Q} is explicitly available; it is just a matter of doing the computations in the correct order. By definition, we have

$$Q_{i,j} = \sum_{k=1}^j L_{i,k}L_{j,k}, \quad i \geq j,$$

where we have used that \mathbf{L} is lower triangular meaning that $L_{i,k} = 0, \forall k > i$. To fix ideas, assume $n = 2$, so that $Q_{1,1} = L_{1,1}^2$, $Q_{2,1} = L_{2,1}L_{1,1}$ and $Q_{2,2} = L_{2,1}L_{2,1} + L_{2,2}^2$. Then we see immediately that we can compute $L_{1,1}$, $L_{2,1}$ and $L_{2,2}$ in this particular order. This generalizes for $n > 2$; we can compute $L_{i,1}$ for $i = 1, \dots, n$ (in this order), then $L_{i,2}$ for $i = 2, \dots, n$, and so on. Due to this simple explicit structure, the Cholesky factorization can be computed quickly". "However, the complexity of the computations is of order $\mathcal{O}(n^3)$."

The natural way to speed up the Cholesky factorization is to make use of a particular structure in the precision matrix or the Cholesky triangle. For GMRFs, the issue is that sparsity in \mathbf{Q} implies (a related) sparsity in \mathbf{L} . The implication is that if we *know* that $L_{j,i} = 0$, then we do not need to compute it. And if the main bulk of \mathbf{L} is zero, then we can achieve great computational savings. In order to understand these issues, we need to understand what \mathbf{L} really means in terms of statistical interpretation. The simulation algorithm for a zero mean GMRFs, which solves $\mathbf{L}^\top \mathbf{x} = \mathbf{z}$, where $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ gives the following result immediately."

Theorem 4.

Let \mathbf{x} be a GMRF wrt the labeled graph \mathcal{G} , mean $\boldsymbol{\mu}$, and a SPD precision matrix \mathbf{Q} . Let \mathbf{L} be the Cholesky triangle of \mathbf{Q} . Then for $i \in \mathcal{V}$,

$$\mathbb{E}(x_i \mid \mathbf{x}_{(i+1):n}) = \mu_i - \frac{1}{L_{i,i}} \sum_{j=i+1}^n L_{j,i}(x_j - \mu_j) \quad \text{and}$$

$$\text{Precision}(x_i \mid \mathbf{x}_{(i+1):n}) = L_{i,i}^2.$$

"Hence the elements of \mathbf{L} have an interpretation as the contribution to the conditional mean and precision for x_i , given all those x_j s where $j > i$. This is in contrast to the elements of \mathbf{Q} , which have a similar interpretation, but where we condition on all other x_j s. A simple consequence of this interpretation is that $Q_{i,i} \geq L_{i,i}^2, \forall i$.

If we merge Equation (3) with Theorem 4, we obtain the following result."

Theorem 5.

Let \mathbf{x} be a GMRF wrt to the labeled graph \mathcal{G} , with mean $\boldsymbol{\mu}$ and SPD precision matrix \mathbf{Q} . Let \mathbf{L} be the Cholesky triangle of \mathbf{Q} and define for $1 \leq i < j \leq n$ the future of i except j as

$$F(i, j) = \{i + 1, \dots, j - 1, j + 1, \dots, n\}.$$

Then

$$x_i \perp x_j \mid \mathbf{x}_{F(i,j)} \iff L_{j,i} = 0.$$

"So, if we consider the marginal distribution of $\mathbf{x}_{i:n}$, then $L_{j,i} = 0$ is equivalent to x_i and x_j being conditionally independent. This is a useful result, as it indicates that if we can determine zeros in \mathbf{L} using conditional independence in the sequence of marginals $\{\mathbf{x}_{i:n}\}$, the marginals $\{\mathbf{x}_{i:n}\}$ are easier to compute through the Cholesky triangle. However, we can use a weaker criterion, which implies that $x_i \perp x_j \mid \mathbf{x}_{F(i,j)}$, the global Markov property in Theorem 3."

Corollary 1.

If $F(i, j)$ separates $i < j$ in \mathcal{G} , then $L_{j,i} = 0$.

"This is the main result. If we can verify that $i < j$ are separated by $F(i, j)$, a operation that only depends on the graph and not the numerical values in \mathbf{Q} , then we know that $L_{j,i} = 0$ no matter what the numerical values in \mathbf{Q} are. If $i < j$ are not separated by $F(i, j)$, then $L_{j,i}$ *can* be zero, but is in general nonzero. Since two neighbors $i \sim j$ are not separated by any set, then $L_{j,i}$ is in general nonzero for neighbors."

Example 5.

"Consider a GMRF with graph as in Figure 1. We then know that $L_{1,1}, L_{2,2}, L_{3,3}, L_{4,4}$ is nonzero, and $L_{2,1}, L_{3,2}, L_{4,3}$, and $L_{4,1}$ is in general nonzero. The two elements remaining are $L_{3,1}$ and $L_{4,2}$, which we check using Corollary 1; nodes 1 and 3 are separated by $F(1, 3) = \{2, 4\}$, so $L_{3,1}$ must be zero, whereas $F(4, 2) = \{3\}$ does not separate 2 and 4 due to node 1, hence $L_{4,2}$ is in general nonzero." A representation of this \mathbf{L} can be seen in (12), with the zero elements represented as empty spaces.

$$\mathbf{L} = \begin{matrix} & \begin{matrix} 1 & 2 & 3 & 4 \end{matrix} \\ \begin{matrix} 1 \\ 2 \\ 3 \\ 4 \end{matrix} & \begin{pmatrix} L_{1,1} & & & \\ L_{2,1} & L_{2,2} & & \\ & L_{3,2} & L_{3,3} & \\ L_{4,1} & L_{4,2} & L_{4,3} & L_{4,4} \end{pmatrix} \end{matrix}. \quad (12)$$

Cholesky Factorization of Band Matrices

"Although only one element of the Cholesky triangle in Example 5 was necessarily zero, we obtain a larger amount of zeros for autoregressive models. Let x be a p th order autoregressive process, $\text{AR}(p)$,

$$x_t \mid x_{t-1}, \dots, x_1 \sim \mathcal{N}(\phi_1 x_{t-1} + \dots + \phi_p x_{t-p}, \sigma^2), \quad t = 1, \dots, n.$$

where we set $x_0 = x_{-1} = \dots = x_{-p+1} = 0$ for simplicity. The precision matrix for an $\text{AR}(p)$ process will then be a band matrix with bandwidth $b_w = p$. Using Corollary 1, it follows immediately that $L_{j,i} = 0$, $j > i$, $\forall j - i > p$, hence \mathbf{L} is a band matrix with the same bandwidth."

Theorem 6.

If \mathbf{Q} is a SPD band matrix with bandwidth b_w , then its Cholesky triangle \mathbf{L} is a lower triangular band matrix with the same bandwidth.

"In this example, only $\mathcal{O}(n(b_w + 1))$ of the $\mathcal{O}(n^2)$ terms in \mathbf{L} are nonzero, which is a significant reduction. A direct consequence is that the algorithm for computing the Cholesky factorization can be simplified; two of the three loops only need to go within the bandwidth, so the complexity is reduced from $\mathcal{O}(n^3)$ to $\mathcal{O}(nb_w^2)$. For fixed b_w , this gives a computational cost, which is linear in n ."

Reordering Techniques: Band Matrices

"The great computational savings we obtained for band matrices naturally raise the question of who then we can use this approach also for "nonband", but sparse, matrices. A rationale for such an approach is that the indices in the graph are arbitrary, hence, we can permute the indices to obtain a small bandwidth, do the computations and perform the inverse permutation on the answer. Formally, let \mathbf{P} be one of the $n!$, $n \times n$ permutation matrices; each row and column of \mathbf{P} has one and only one nonzero entry, which is 1. The transpose of a permutation matrix is the inverse permutation, $\mathbf{P}^\top \mathbf{P} = \mathbf{I}$. For example, $\mathbf{Q}\boldsymbol{\mu} = \mathbf{b}$ can be solved as follows; multiply both sides with \mathbf{P}

$$\underbrace{(\mathbf{P}\mathbf{Q}\mathbf{P}^\top)}_{\tilde{\mathbf{Q}}} \underbrace{\mathbf{P}\boldsymbol{\mu}}_{\tilde{\boldsymbol{\mu}}} = \underbrace{\mathbf{P}\mathbf{b}}_{\tilde{\mathbf{b}}}$$

solve $\tilde{\mathbf{Q}}\tilde{\boldsymbol{\mu}} = \tilde{\mathbf{b}}$, and then apply the inverse permutation to obtain the solution $\boldsymbol{\mu} = \mathbf{P}^\top \tilde{\boldsymbol{\mu}}$.

The next issue is how to permute the sparse matrix in order to obtain a (possible) small bandwidth. This issue is somewhat more technical, but there is a huge literature in computer science, and good working algorithms. So any algorithm that runs quick and gives reasonable results is fine. Figure 4 displays an example, where in panel (a) we display the precision matrix found from a spatial application

- (SECTION 4.2.2 OF GMRFs BOOK),

in (b) the reordered precision matrix using the Gibbs-Poole-Stockmeyer reordering algorithm




Figure 4: Panel (a) displays the nonzeros of the 380×380 precision matrix, panel (b) displays the reordered precision matrix with bandwidth 38, and panel (c) displays the Cholesky triangle of the band matrix (b).

- (J.G. Lewis. Algorithm 582: The Gibbs-Poole-Stockmeyer and Gibbs-King algorithms for reordering sparse matrices. *ACM Transactions on Mathematical Software*, 8(2):190-194, June 1982.),

and in (c), the Cholesky triangle of the reordered precision matrix. The bandwidth after reordering is 38."

Reordering Techniques: General Sparse Matrices